

Reflective Report

about Introduction to Natural Language

Processing Concepts Exercise

Iverson Labastida
BSIT
Jose Rizal University
Mandaluyong, Philippines
labastidaiverson@gmail.com

Abstract—This exercise introduces me to foundational concepts in natural language processing like the techniques use in analyzing text data, concepts like tokenization, statistical analysis like using frequency of words that appeared in a corpus, also about semantic language models on which it introduces me the concept of embeddings and after that I did the hands-on exercise provided by the course, and I also did an assessment to test my knowledge.

I. EXPERIENTIAL & PRACTICAL INSIGHTS

The exercise demonstrated that generative and specialized language tools can effectively illustrate core NLP workflows summarization, language detection, and PII extraction while exposing practical tradeoffs between local and cloud deployments; the Phi 3.5-mini model produced coherent, concise summaries suitable for low-bandwidth scenarios but occasionally missed nuanced content, the Language Playground reliably identified supported languages for clear inputs yet struggled with short or mixed-language labels, and the PII analyzer detected obvious personal data using a hybrid of generative cues and pattern matching but exhibited false positives and negatives on nonstandard formats; overall, the lab highlights the pedagogical value of interactive, editable interfaces and local execution for privacy and experimentation, while underscoring that production-grade analytics require larger models or cloud services, deterministic validation (regex/rule engines), and human review to mitigate errors, bias, and privacy risks.

II. CORE LEARNING OUTCOMES

I found that tokenization choices word, sub-word, or character require trade-offs between vocabulary size, handling rare or out-of-vocabulary tokens, and computational cost, and I prefer sub-word tokenization for its practical balance of efficiency; I also see TF-IDF as a fast, interpretable, sparse baseline that captures term importance but ignores semantics, while embeddings offer dense, semantic representations that enable similarity search and algebraic reasoning; finally, I value context vectors from transformer-style attention for refining token meanings based on surrounding text, which improves disambiguation and downstream tasks, so my rule of thumb is: use TF-IDF for lightweight, explainable baselines, static embeddings for semantic retrieval, and contextual embeddings when nuance and contextual understanding matter.

III. COURSE APPLICATION & IMPACT

The introductory concepts of text analysis and natural language processing (NLP) provide me with a strong foundation for the advanced topics that I will encounter throughout this course. By understanding tokenization, text normalization, statistical methods such as TF-IDF, semantic language models, and vector embeddings, I am better prepared to complete upcoming laboratory exercises involving text processing, classification, sentiment analysis, and information extraction. These concepts also equip me with the analytical and technical skills needed to design and implement my capstone project, where I can develop intelligent applications that process unstructured text, generate meaningful insights, and solve real-world problems using modern AI and NLP techniques. Overall, mastering these fundamentals will help me confidently apply both traditional and semantic approaches to build accurate, efficient, and practical AI-powered solutions.

IV. CONCLUSION & FUTURE WORK

In conclusion, this study has strengthened my understanding of the fundamental concepts of text analysis, natural language processing (NLP), and semantic language models that form the foundation of modern artificial intelligence applications. I have learned how traditional techniques such as tokenization, frequency analysis, and TF-IDF complement advanced approaches based on embeddings and large language models (LLMs) to extract meaningful information from unstructured text. These insights have enhanced my appreciation of the capabilities and limitations of current NLP technologies while preparing me for more advanced coursework. As I continue my studies, I plan to deepen my knowledge of transformer architectures, prompt engineering, retrieval-augmented generation (RAG), fine-tuning techniques, and responsible AI practices to develop more accurate, efficient, and ethical intelligent systems for real-world applications.